SmolVLM for Dense Video Captioning

Week 4: Internship Update

- Adaptive Thresholding in PySceneDetect
  - tested the default threshold and min_scene_len for Contentdetector and they do not work well always because _calculate_frame_score returns 0.0 for the first frame and it needs at least two frames to detect a shot which might be a problem
  - still used ContectDetector but with observing all the differences and setting the threshold in two ways:
    - threshold = mean + 2std: in the Gaussian case 95.4% of data
    - threshold = 98th percentile of differences: worked okay and extracted at least two frames

- Dynamic Frame Sampling for Multimodal Large Language Model Video Understanding

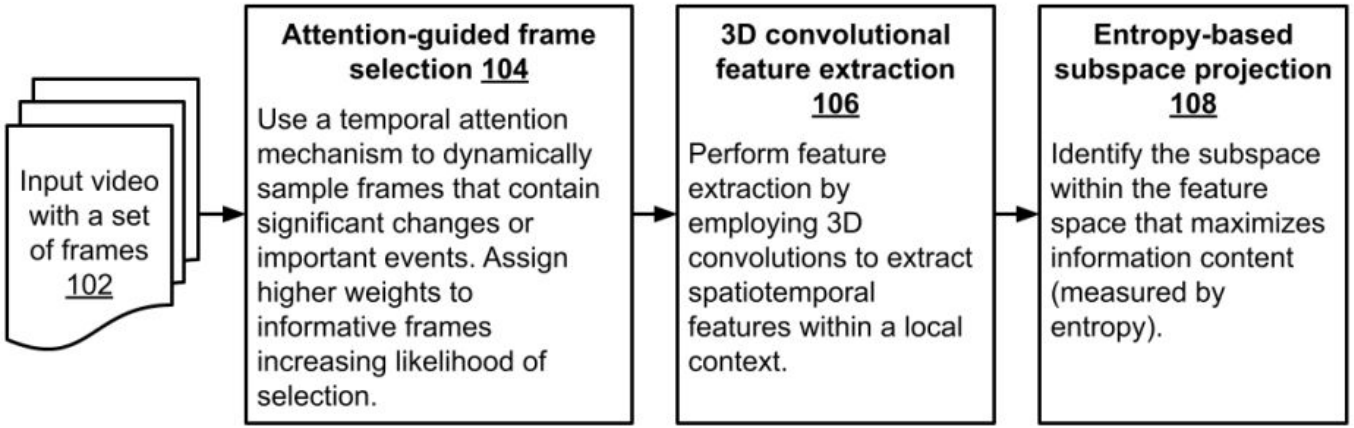

Fig. 1: Dynamic frame sampling for multimodal large language model (LLM) video understanding

- DFS-QA: Dynamic Frame Selection for Better Video Question Answering
    - mostly for the question answering task, not very applicable

- A New Efficient Hybrid Technique for Human Action Recognition Using 2D Conv-RBM and LSTM with Optimized Frame Selection
    - Human Action Recognition task
    - SMART frame selection
    - two-dimensional convolutional restricted Boltzmann machine for spatial feature extraction
    - LSTM for temporal modeling
    - grayscale images
    - SMART Frame Selection for Action Recognition

- SMART Frame Selection for Action Recognition
  - CNN: MobileNet trained on ImageNet + averaged representation of 10 most likely classes (text embedded with GloVe)
  - Single-Frame Selector: MLP BUT I would need classes
  - Global Selector: Attention model over the entire video
  - Multiplication of weights gives the final weight

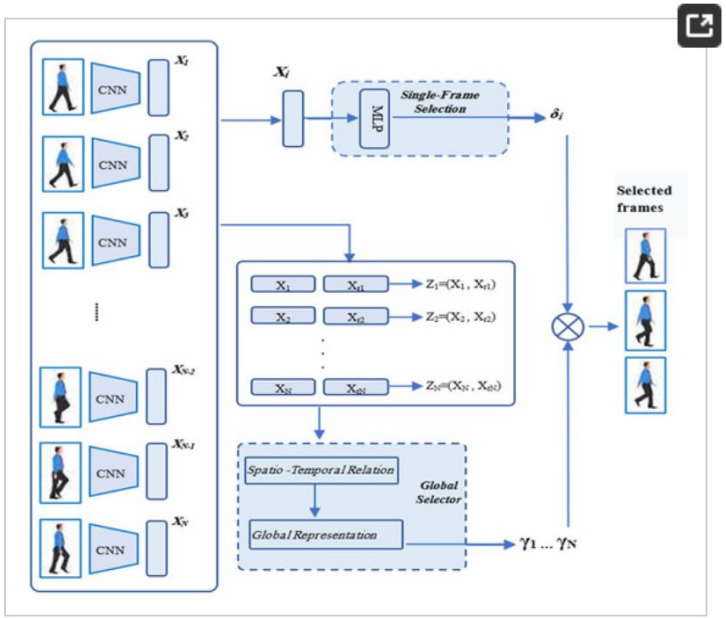

- MaxInfo: A Training-Free Key-Frame Selection Method Using Maximum Volume for Enhanced Video Understanding

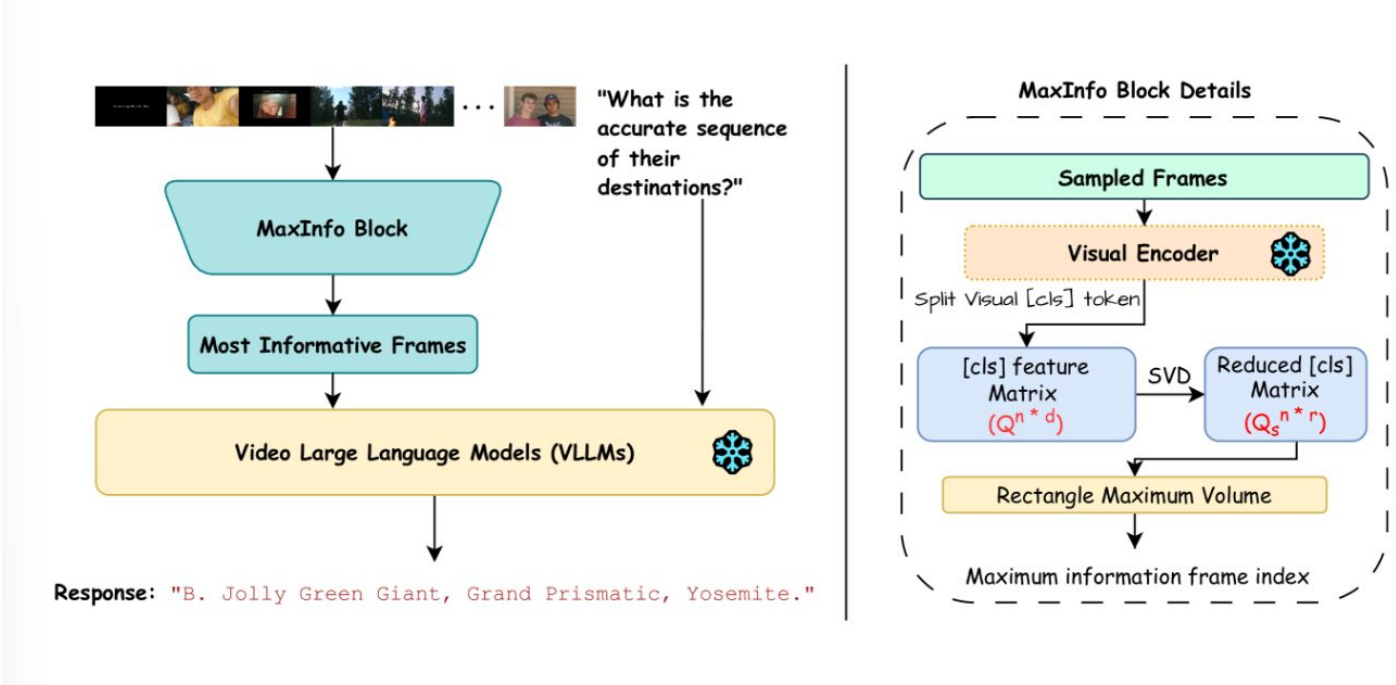

- MaxInfo: A Training-Free Key-Frame Selection Method Using Maximum Volume for Enhanced Video Understanding

Algorithm 1 MaxInfo Block: SVD + MaxVol for Keyframe Selection

Input: A set of n frames $\mathbf{I} = \{i_1, i_2, \ldots, i_n\}$

Embedding: Convert each frame i_j into a [CLS] embedding:

$$q_n = \text{flatten}\big(\text{clip\_model}(i_n)\big), \mathbf{Q} = \begin{bmatrix} q_1 \\ q_2 \\ \vdots \\ q_n \end{bmatrix} \in \mathbb{R}^{n \times d}.$$

SVD Reduction: Perform truncated SVD on $\mathbf{Q}$:

$$\mathbf{Q} \approx \mathbf{U}_r \mathbf{\Sigma}_r \mathbf{V}_r^{\top} \quad \rightarrow \quad \mathbf{Q_s} = \mathbf{U}_r \in \mathbb{R}^{n \times r}.$$

MaxVol Selection: Run $\text{rect\_maxvol}(\mathbf{Q_s}, \text{tol})$ to find pivot indices:

$$\text{piv} = \text{rect\_maxvol}(\mathbf{Q_s}, \text{Tol}),$$

identifying rows (frames) that span the reduced embedding space.

Output: Indices $\mathbf{piv}$ of the most informative keyframes.

- MaxInfo: A Training-Free Key-Frame Selection Method Using Maximum Volume for Enhanced Video Understanding
    - Truncated SVD(Singular Value Decomposition)
        - Reducing feature dimension
        - Optimal s could be calculated
    - Rectangular MaxVolume Selection
    - Scene-Aware MaxInfo

$$\text{rect-vol}(\mathbf{A}) = \sqrt{\det(\mathbf{A}\,\mathbf{A}^T)} \qquad \mathbf{r} = \arg\max_{\mathbf{r}} \text{rect-vol}\left(\mathbf{Q}_s(\mathbf{r}, :)\right)$$

- MaxInfo: A Training-Free Key-Frame Selection Method Using Maximum Volume for Enhanced Video Understanding

Model	Size	VideoMME (wo/w-subs)	EgoSchema	LongVideoBench
LLaVA-Video (Zhang et al., 2024c)	7B	$63.3/69.7_{(64)}$	$57.3_{(64)}$	$58.2_{(64)}$
+ MaxInfo	7B	$64.2/71.4_{64\rightarrow(6,64)}$	$63.7_{128\rightarrow(64,64)}$	$61.5_{128\rightarrow(1,64)}$
△		+0.9%/+1.7%	+6.4%	+3.3%
LLaVA- (Zhang et al., 2024c)	72B	$70.5/76.9_{(64)}$	$65.6_{(64)}$	$61.9_{(64)}$
+ MaxInfo	72B	$70.2/77.6_{64\rightarrow(6,64)}$	$69.4_{128\rightarrow(64,64)}$	$64.9_{128\rightarrow(64,1)}$
△		-0.3%/+0.7%	+3.8%	+3%
Qwen2-VL (Wang et al., 2024a)	2B	$55.6/60.4_{(786)}$	$54.9_{(180)}$	$47.3_{(256)}$
+ MaxInfo	2B	$57.0/61.6_{256\rightarrow(254,4)}$	$57.2_{180\rightarrow(180,12)}$	$48.8_{256\rightarrow(224,1)}$
△		+1.4%/+1.2%	+2.3%	+1.5%
Qwen2-VL (Wang et al., 2024a)	7B	$63.3/69.0_{(768)}$	$66.7_{(180)}$	$53.7_{(256)}$
+ MaxInfo	7B	$62.1/70.0_{256\rightarrow(254,4)}$	$64.3_{180\rightarrow(180,12)}$	$55.7_{256\rightarrow(224,1)}$
△		-1.2/+1.0%	+2.4%	+2.0%

Model	Size	Visual encoder	Param.	Acc.
LLaVA-Video	7B	CLIP-ViT-Large	427.9 M.	58.94
LLaVA-Video	7B	CLIP-ViT-Base	149.6 M.	58.79
LLaVA-Video	7B	DINO2-Base	86.6 M.	58.94
LLaVA-Video	7B	DINO2-Large	304.4 M.	58.86
LLaVA-Video	7B	SigLIP-Base-224	203.2 M.	**59.76**
LLaVA-Video	7B	SigLIP-Base-384	878 M.	59.24

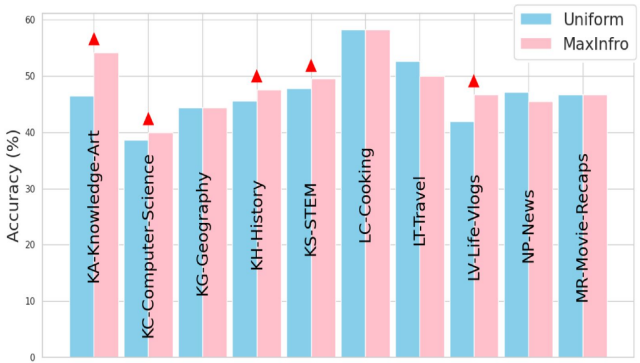

Figure 3: Accuracy comparison between Uniform Sampling and MaxInfo on Video-MME (Qwen2-VL-2B).